

Clarinet: WAN-Aware Optimization for Analytics Queries

Raajay Viswanathan, Ganesh Ananthanarayanan, Aditya Akella



Microsoft®
Research

Overview



- Web apps hosted on multiple DCs → Low latency access to end-user

Overview





- Web apps hosted on multiple DCs → Low latency access to end-user

Overview




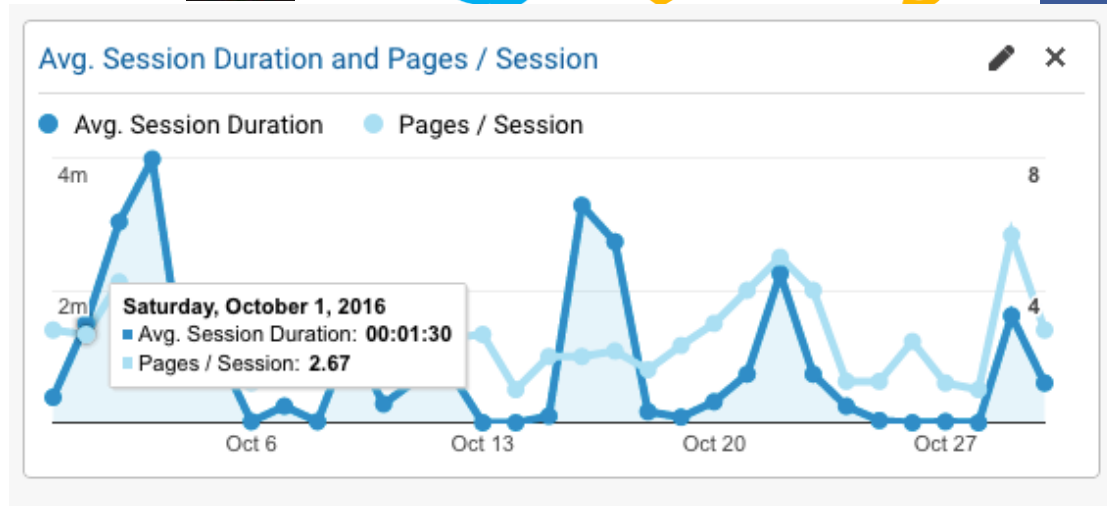
- Web apps hosted on multiple DCs → Low latency access to end-user

Overview



- Web apps hosted on multiple DCs → Low latency access to end-user

Overview





- Web apps hosted on multiple DCs → Low latency access to end-user
- Need efficient methods to analyze data located in **multiple** data centers

Centralized Aggregation is Wasteful



Centralized Aggregation is Wasteful

```
SELECT * ... FROM .. WHERE .. ;
```





Centralized Aggregation is Wasteful

```
SELECT * ... FROM .. WHERE .. ;
```



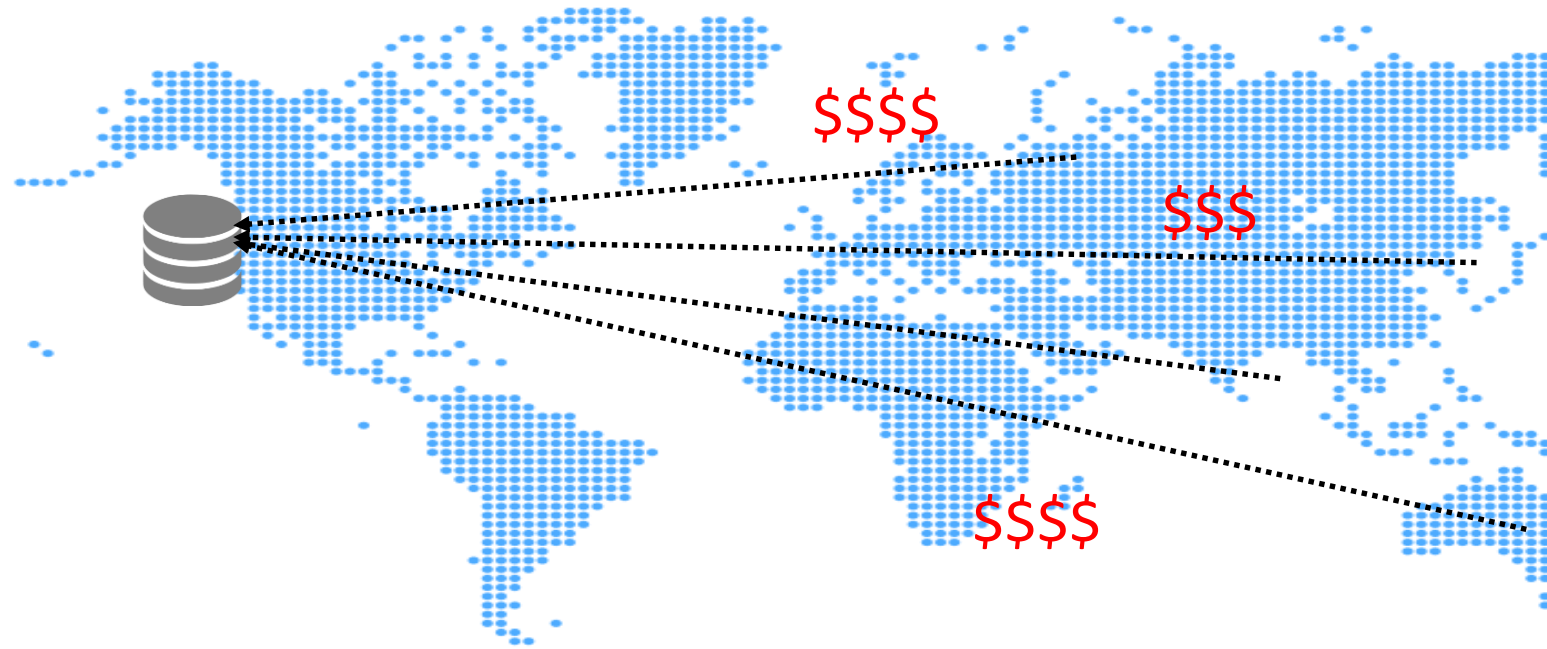

Intra-data center
Analytics
Framework



- Available WAN bandwidth is limited → Aggregation latency overhead

Centralized Aggregation is Wasteful

```
SELECT * ... FROM .. WHERE .. ;
```

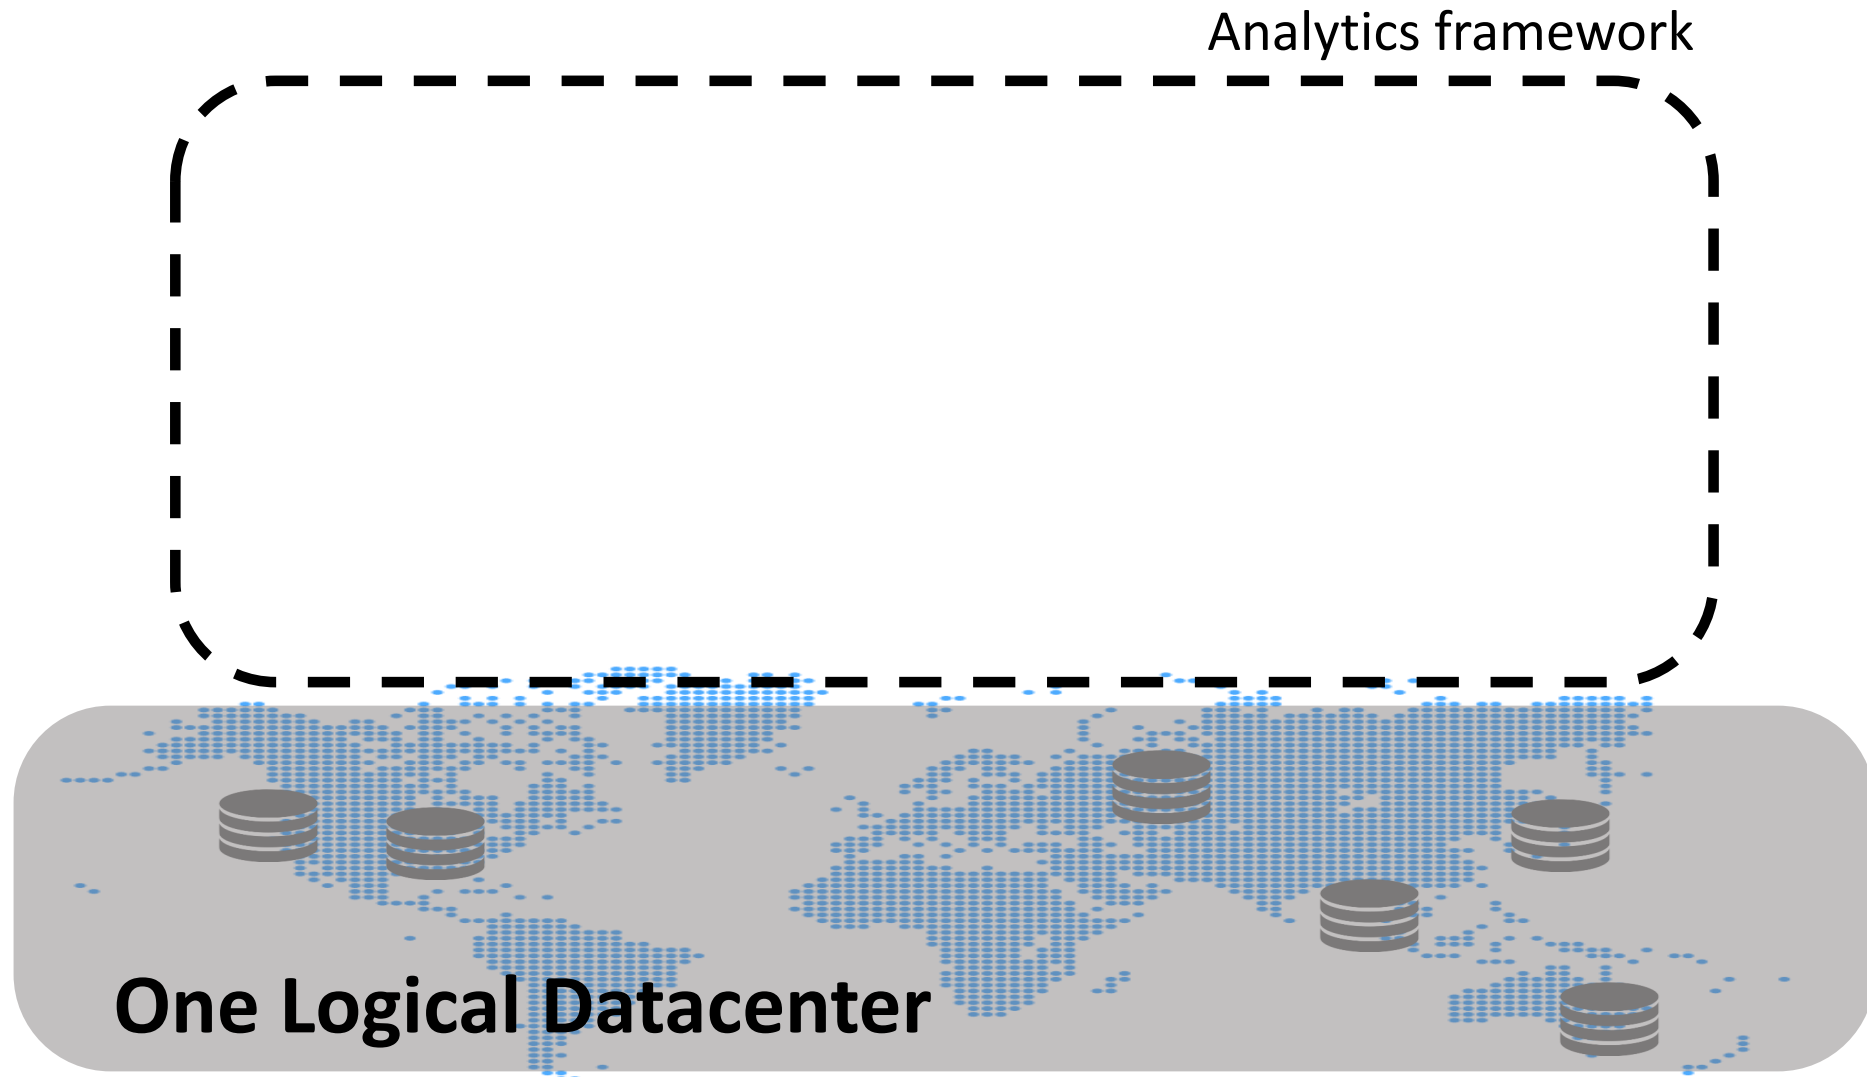


- Available WAN bandwidth is limited → Aggregation latency overhead
- WAN links are expensive → High data transfer cost


Geo-distributed Analytics




Geo-distributed Analytics



Geo-distributed Analytics



Geo-distributed Analytics



Geo-distributed Analytics

SELECT * ... FROM .. WHERE .. ;

Geo-distributed Analytics framework

Query Optimizer

Multi-stage parallelizable jobs


Distributed Execution Layer

Distributed Storage Layer

One Logical Datacenter

Requires
WAN-aware
optimization

Geo-distributed Analytics



Geo-distributed Analytics

SELECT * ... FROM .. WHERE .. ;

Geo-distributed Analytics framework

Query Optimizer

Multi-stage parallelizable jobs

Distributed Execution Layer

Distributed Storage Layer

Clarinet

Requires
WAN-aware
optimization

One Logical Datacenter

Geo-distributed Analytics

SELECT * ... FROM .. WHERE .. ;

Geo-distributed Analytics framework

Query Optimizer

Multi-stage parallelizable jobs

Distributed Execution Layer

Distributed Storage Layer

2.7x reduction in
query runtime
Clarinet

Requires
WAN-aware
optimization


One Logical Datacenter

WAN Aware Query Optimization




T₁, T₂, T₃: Tables storing click logs

WAN Aware Query Optimization



T₁, T₂, T₃: Tables storing click logs

WAN Aware Query Optimization




T1, T2, T3: Tables storing click logs

QUERY

```
SELECT T1.user, T1.latency, T2.latency, T3.latency
FROM   T1, T2, T3
WHERE  T1.user == T2.user AND T1.user == T3.user
AND    T1.device == T2.device == T3.device == "mobile";
```


WAN Aware Query Optimization




T1, T2, T3: Tables storing click logs

QUERY

```
SELECT T1.user, T1.latency, T2.latency, T3.latency
FROM T1, T2, T3
WHERE T1.user == T2.user AND T1.user == T3.user
AND T1.device == T2.device == T3.device == "mobile";
```




WAN Aware Query Optimization




T1, T2, T3: Tables storing click logs

QUERY

```
SELECT T1.user, T1.latency, T2.latency, T3.latency
FROM T1, T2, T3
WHERE T1.user == T2.user AND T1.user == T3.user
AND T1.device == T2.device == T3.device == "mobile";
```




WAN Aware Query Optimization




T1, T2, T3: Tables storing click logs

QUERY

```
SELECT T1.user, T1.latency, T2.latency, T3.latency
FROM T1, T2, T3
WHERE T1.user == T2.user AND T1.user == T3.user
AND T1.device == T2.device == T3.device == "mobile";
```



WAN Aware Query Optimization




T1, T2, T3: Tables storing click logs


QUERY

```
SELECT T1.user, T1.latency, T2.latency, T3.latency
FROM T1, T2, T3
WHERE T1.user == T2.user AND T1.user == T3.user
AND T1.device == T2.device == T3.device == "mobile";
```

Plan running time: **41 s**



WAN Aware Query Optimization




T1, T2, T3: Tables storing click logs


QUERY

```
SELECT T1.user, T1.latency, T2.latency, T3.latency
FROM T1, T2, T3
WHERE T1.user == T2.user AND T1.user == T3.user
AND T1.device == T2.device == T3.device == "mobile";
```

Plan running time: 41 s



WAN Aware Query Optimization




T1, T2, T3: Tables storing click logs


QUERY

```
SELECT T1.user, T1.latency, T2.latency, T3.latency
FROM T1, T2, T3
WHERE T1.user == T2.user AND T1.user == T3.user
AND T1.device == T2.device == T3.device == "mobile";
```

Plan running time: 41 s



WAN Aware Query Optimization




T1, T2, T3: Tables storing click logs

QUERY


```
SELECT T1.user, T1.latency, T2.latency, T3.latency  
FROM T1, T2, T3  
WHERE T1.user == T2.user AND T1.user == T3.user  
AND T1.device == T2.device == T3.device == "mobile";
```

Plan running time: 41 s




Plan A

Plan running time: 20.96 s




Plan B

Plan running time: 17.6 s



Plan C

WAN Aware Query Optimization




T1, T2, T3: Tables storing click logs


QUERY

```
SELECT T1.user, T1.latency, T2.latency, T3.latency
FROM T1, T2, T3
WHERE T1.user == T2.user AND T1.user == T3.user
AND T1.device == T2.device == T3.device == "mobile";
```


Plan running time: 41 s




Plan running time: 20.96 s



Plan running time: 17.6 s



WAN Aware Query Optimization

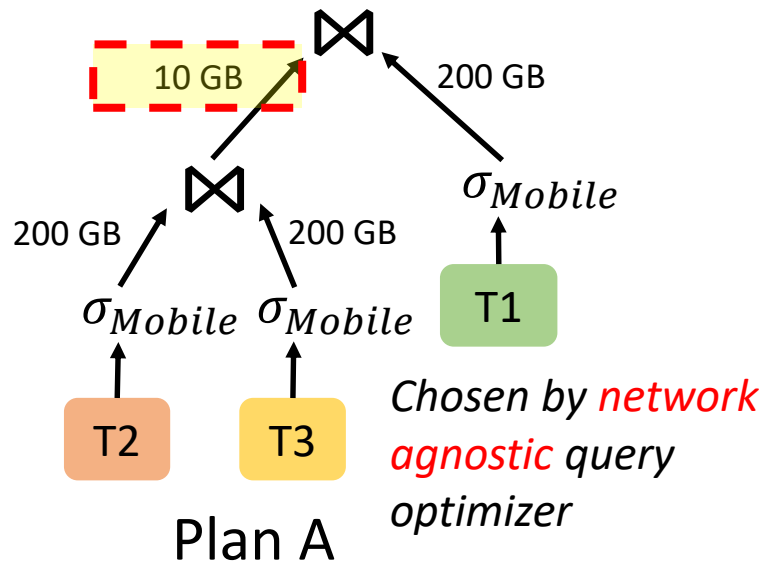


T1, T2, T3: Tables storing click logs


QUERY

```
SELECT T1.user, T1.latency, T2.latency, T3.latency
FROM T1, T2, T3
WHERE T1.user == T2.user AND T1.user == T3.user
AND T1.device == T2.device == T3.device == "mobile";
```


Plan running time: 41 s




Plan running time: 20.96 s



Plan running time: 17.6 s



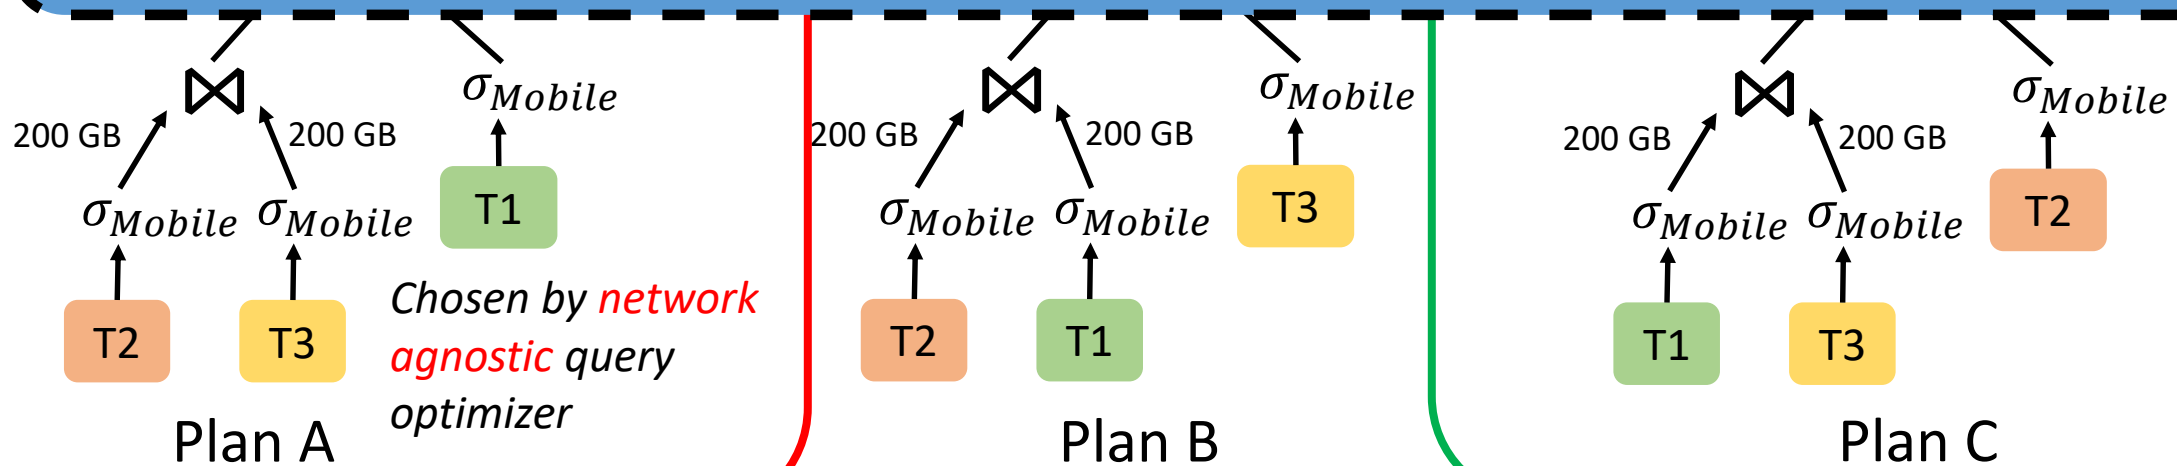
WAN Aware Query Optimization



QUERY

```
SELECT T1.user, T1.latency, T2.latency, T3.latency  
FROM T1, T2, T3  
WHERE T1.user == T2.user AND T1.user == T3.user  
AND T1.device == T2.device == T3.device == "mobile";
```


WAN-aware query optimizer that uses network transfer duration to choose query plans





Outline

1. Motivation
2. Challenges in choosing query plan based on WAN transfer durations
3. Solution
 - Single query
 - Multiple simultaneous queries
4. Experimental Evaluation




Other factors also affect query plan run time



Other factors also affect query plan run time

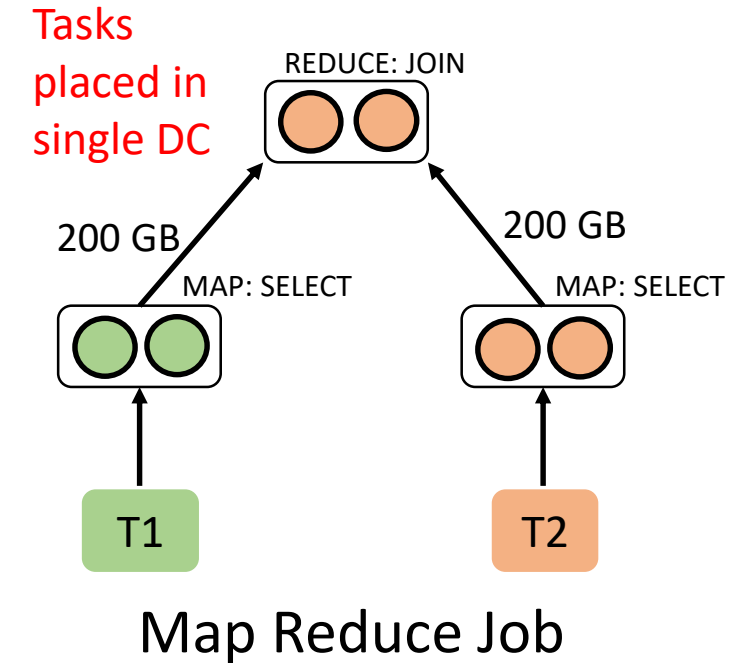




Other factors also affect query plan run time

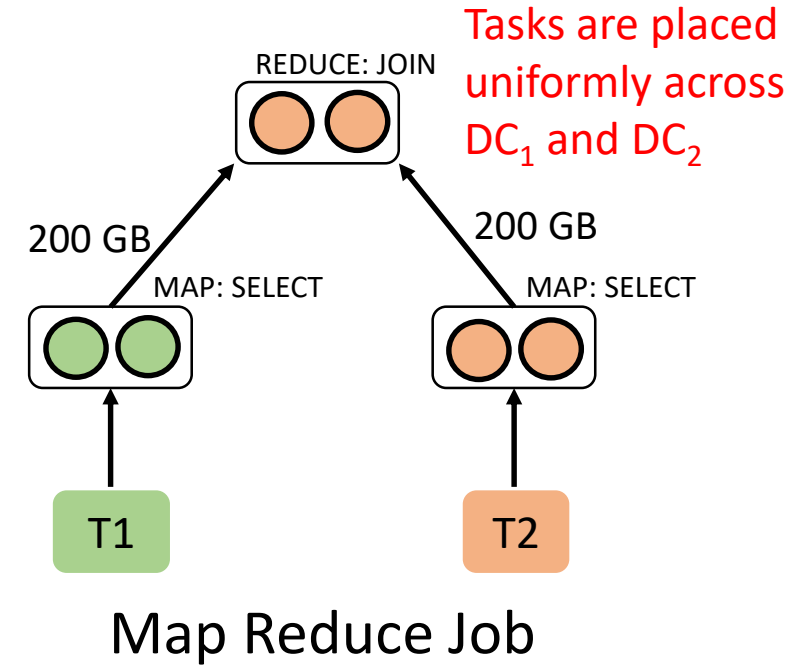




Map Reduce Job

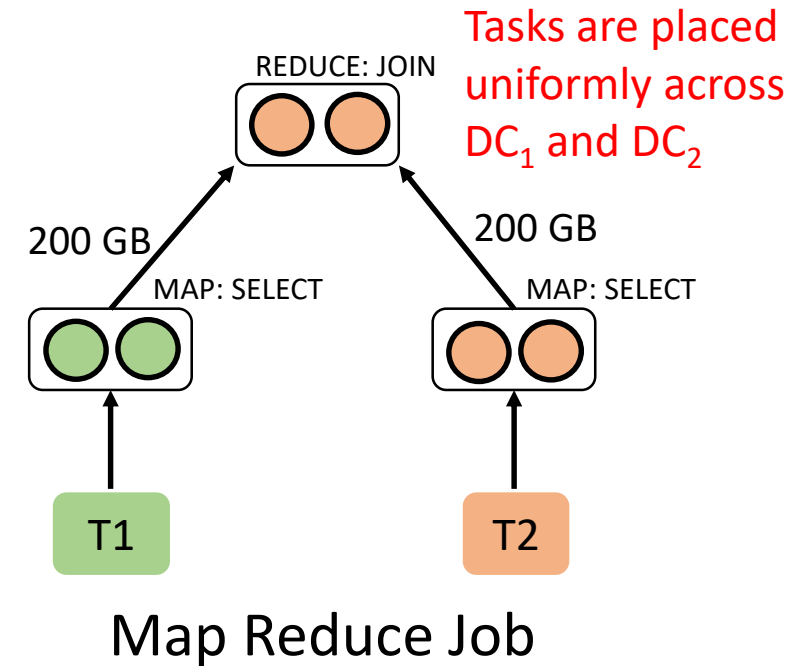


Other factors also affect query plan run time



Other factors also affect query plan run time



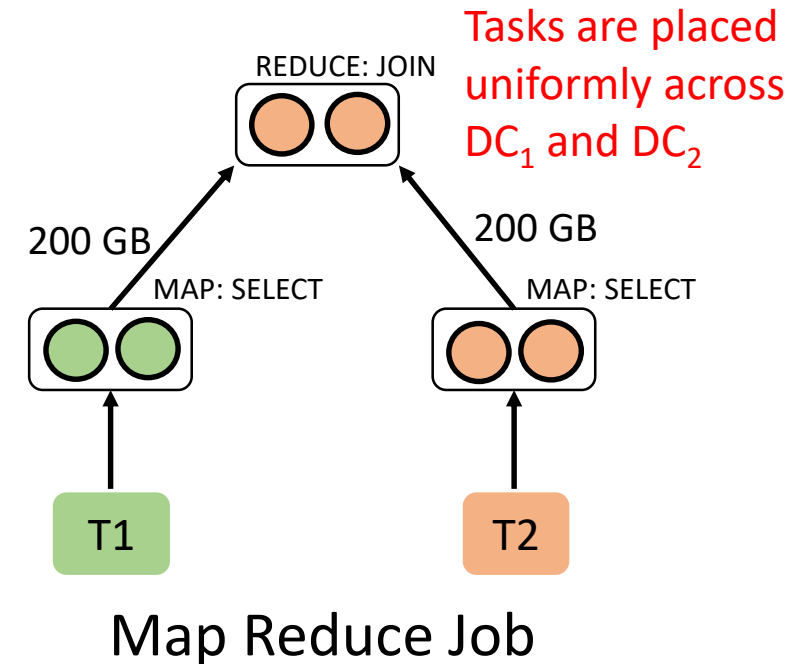


Other factors also affect query plan run time



While evaluating different query plans

1. Plan A: 41 s
2. Plan B: 20.96
3. Plan C: 17.6 s




Other factors also affect query plan run time



While evaluating different query plans

1. Plan A: ~~41 s~~ 20.5 s
2. Plan B: ~~20.96~~ 11.2 s
3. Plan C: 17.6 s




Other factors also affect query plan run time



While evaluating different query plans

1. Plan A: ~~41 s~~ 20.5 s
2. Plan B: ~~20.96~~ 11.2 s
3. Plan C: 17.6 s




Other factors also affect query plan run time



While evaluating different query plans

1. Plan A: ~~41 s~~ 20.5 s
2. Plan B: ~~20.96~~ 11.2 s
3. Plan C: 17.6 s

Other factors also affect query plan run time




Choose query plan based on:

1. Best available task placements

1. Plan A: ~~41 s~~ 20.5 s
2. Plan B: ~~20.96~~ 11.2 s
3. Plan C: 17.6 s

Other factors also affect query plan run time



Choose query plan based on:


1. Best available task placements
2. Schedule of network transfers

- WI
1. Plan A: ~~41 s~~ 20.5 s
 2. Plan B: ~~20.96~~ 11.2 s
 3. Plan C: 17.6 s

Joint plan selection, placement and scheduling


Joint plan selection, placement and scheduling

```
SELECT * FROM ... WHERE... ;
```



Multiple query plans (join orders) per query


Joint plan selection, placement and scheduling



Multiple query plans (join orders) per query

Assign parallelism for each stage

Joint plan selection, placement and scheduling



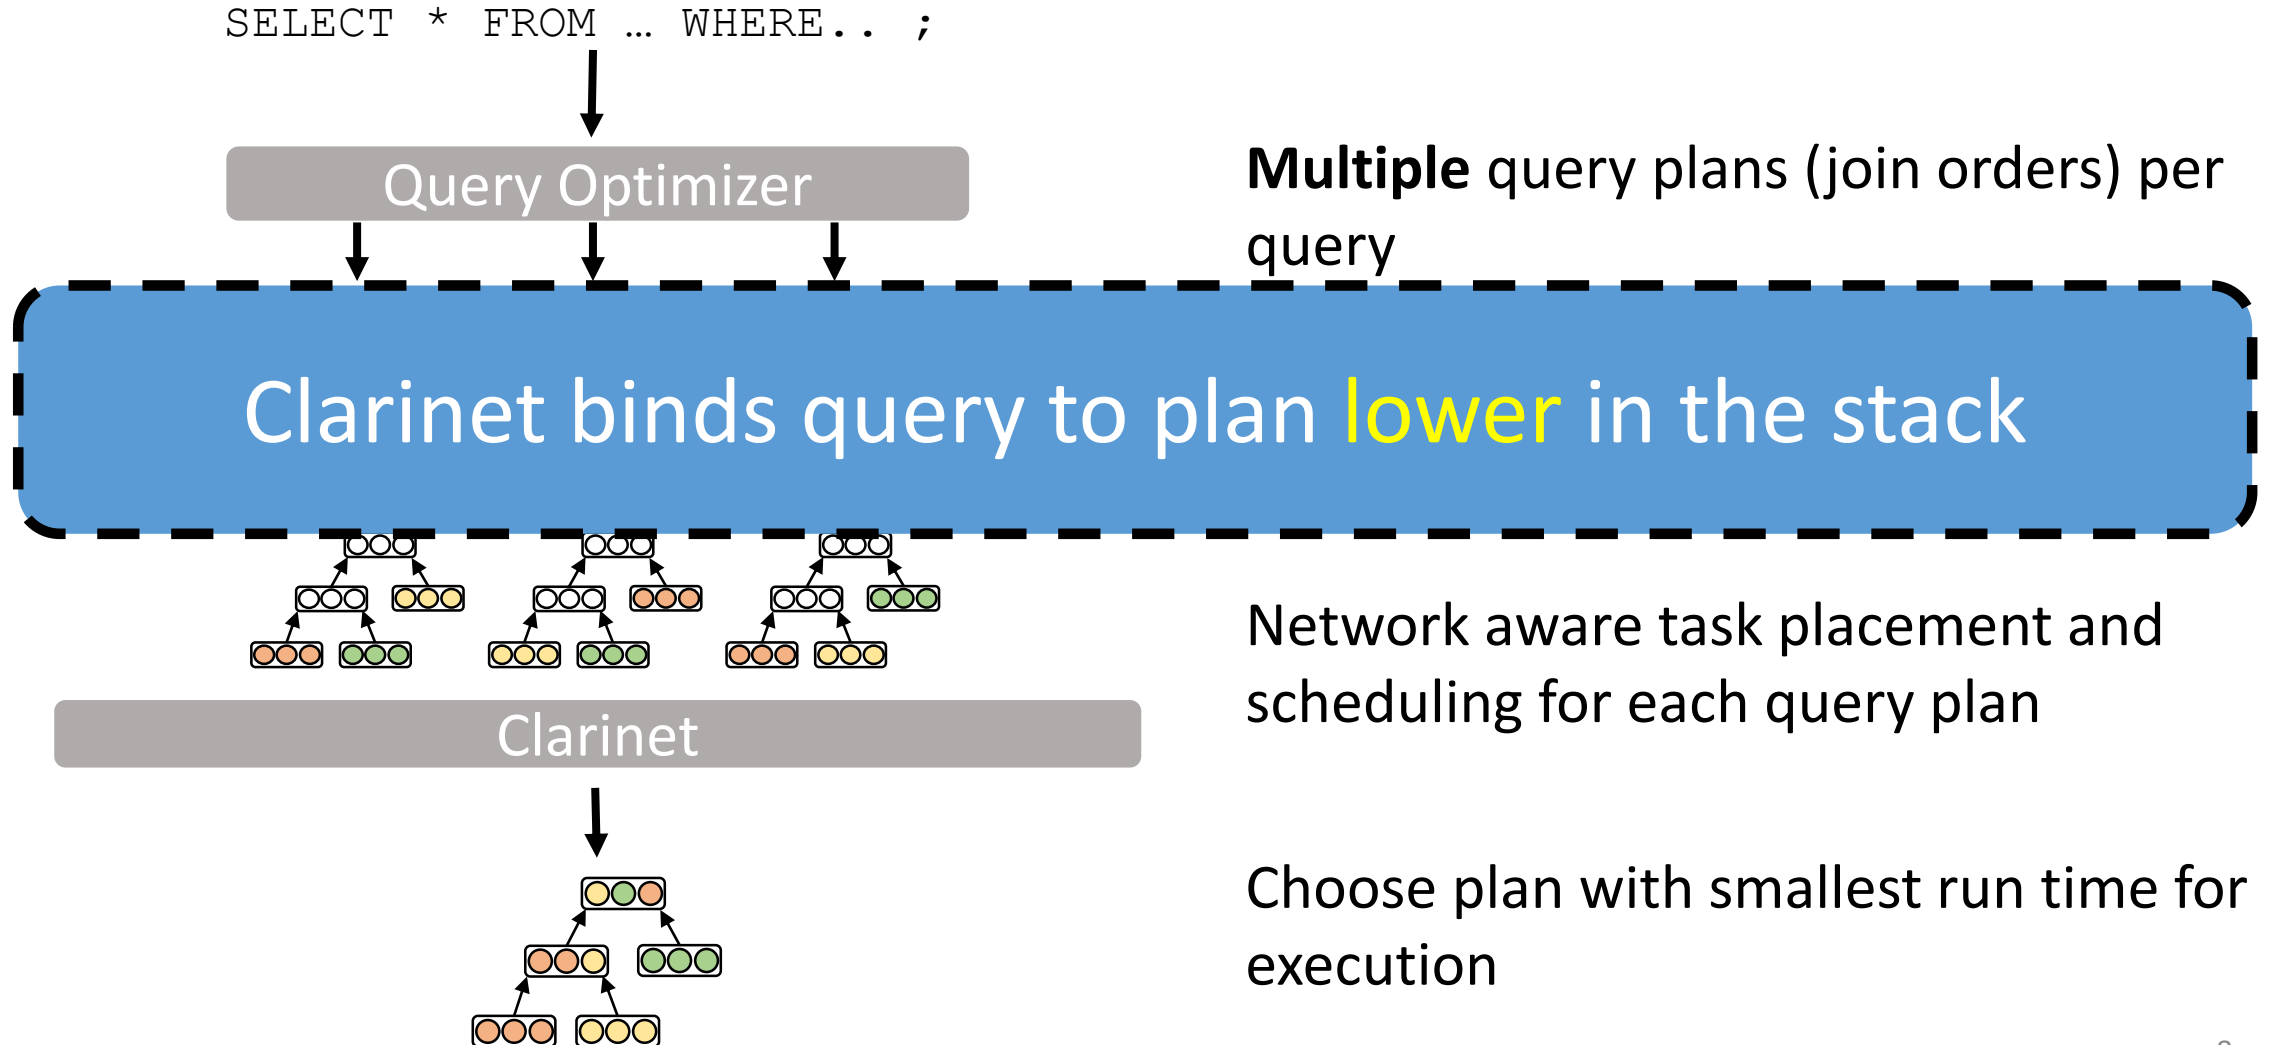
Multiple query plans (join orders) per query

Assign parallelism for each stage


Network aware task placement and scheduling for each query plan

Choose plan with smallest run time for execution

Joint plan selection, placement and scheduling




Network aware placement and scheduling




Network aware placement and scheduling

- Task placement decided greedily one stage at a time
 - Minimize per stage run time




Network aware placement and scheduling

- Task placement decided greedily one stage at a time
 - Minimize per stage run time
- Scheduling of network transfers
 - Determines start times of inter-DC network transfers




Network aware placement and scheduling

- Task placement decided greedily one stage at a time
 - Minimize per stage run time
- Scheduling of network transfers
 - Determines start times of inter-DC network transfers
 - Formulate a Binary Integer Linear Program to solve scheduling
 - Factors transfer dependencies




How to extend the late-binding strategy to multiple queries?

Queries affect each others' run time



Queries affect each others' run time




QUERY 1

```
SELECT ...  
device == "mobile"  
...;
```

QUERY 2

```
SELECT ...  
genre == "pc"  
...;
```

Queries affect each others' run time





QUERY 1

```
SELECT ...  
device == "mobile"  
...;
```


QUERY 2

```
SELECT ...  
genre == "pc"  
...;
```



Same query plan (Plan C) for Query 1 and Query 2

Queries affect each others' run time





QUERY 1

```
SELECT ...  
device == "mobile"  
...;
```


QUERY 2

```
SELECT ...  
genre == "pc"  
...;
```



Same query plan (Plan C) for Query 1 and Query 2

Queries affect each others' run time





Different query plans for Query 1 (Plan C) and Query 2 (Plan B)

QUERY 1


```
SELECT ...  
device == "mobile"  
...;
```

QUERY 2

```
SELECT ...  
genre == "pc"  
...;
```



Queries affect each others' run time



No contention of network links



Different query plans for Query 1 (Plan C)
and Query 2 (Plan B)

QUERY 1


```
SELECT ...  
device == "mobile"  
...;
```

QUERY 2

```
SELECT ...  
genre == "pc"  
...;
```



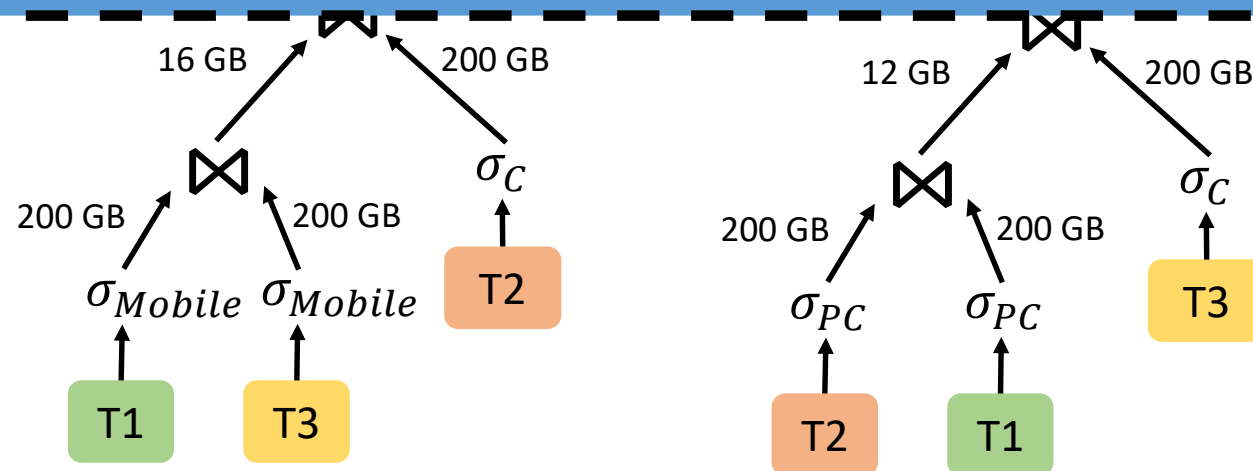
Queries affect each others' run time




Choosing execution plans **jointly** for multiple queries improves performance

QUERY 1
SELECT ...
device == "mobile"
...;

QUERY 2
SELECT ...
genre == "pc"
...;




Iterative Shortest Job First




- Best combination → minimize average completion
 - Computationally intractable

Iterative Shortest Job First



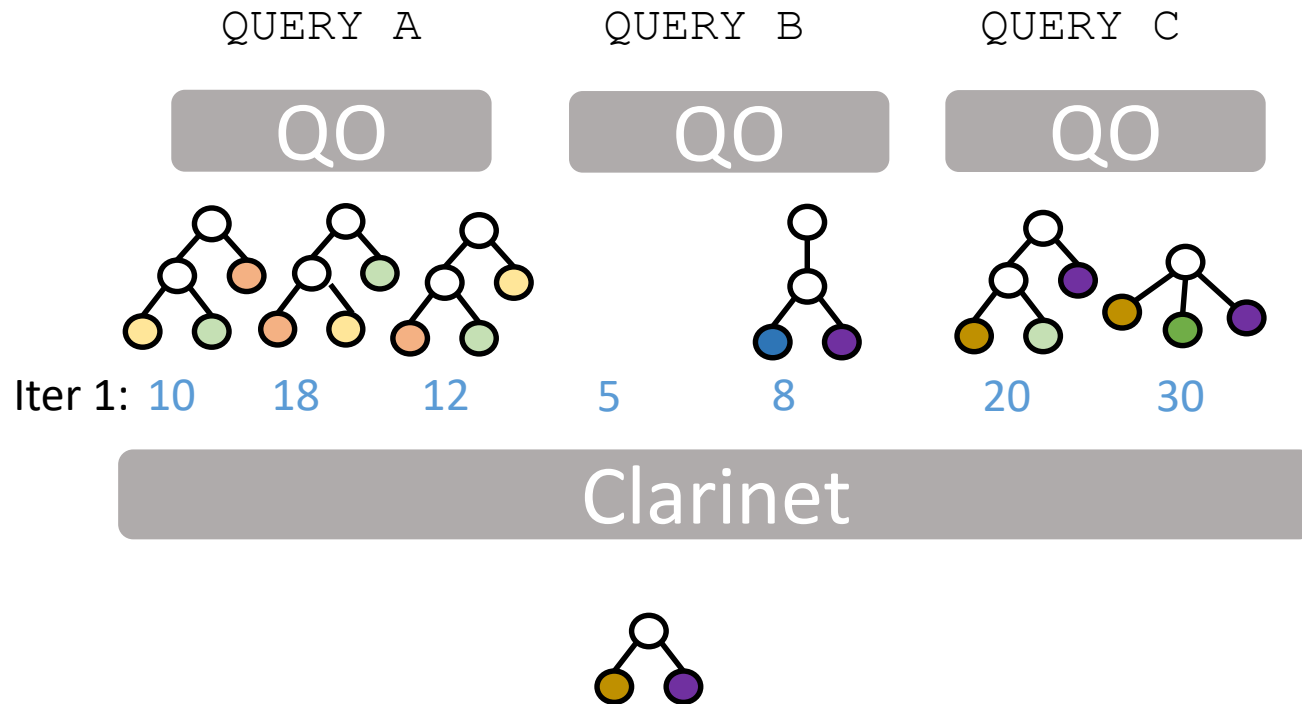
- Best combination → minimize average completion
 - Computationally intractable
- Iterative Shortest Job First (SJF) scheduling heuristic
 1. Pick shortest physical query plan in each iteration

Iterative Shortest Job First



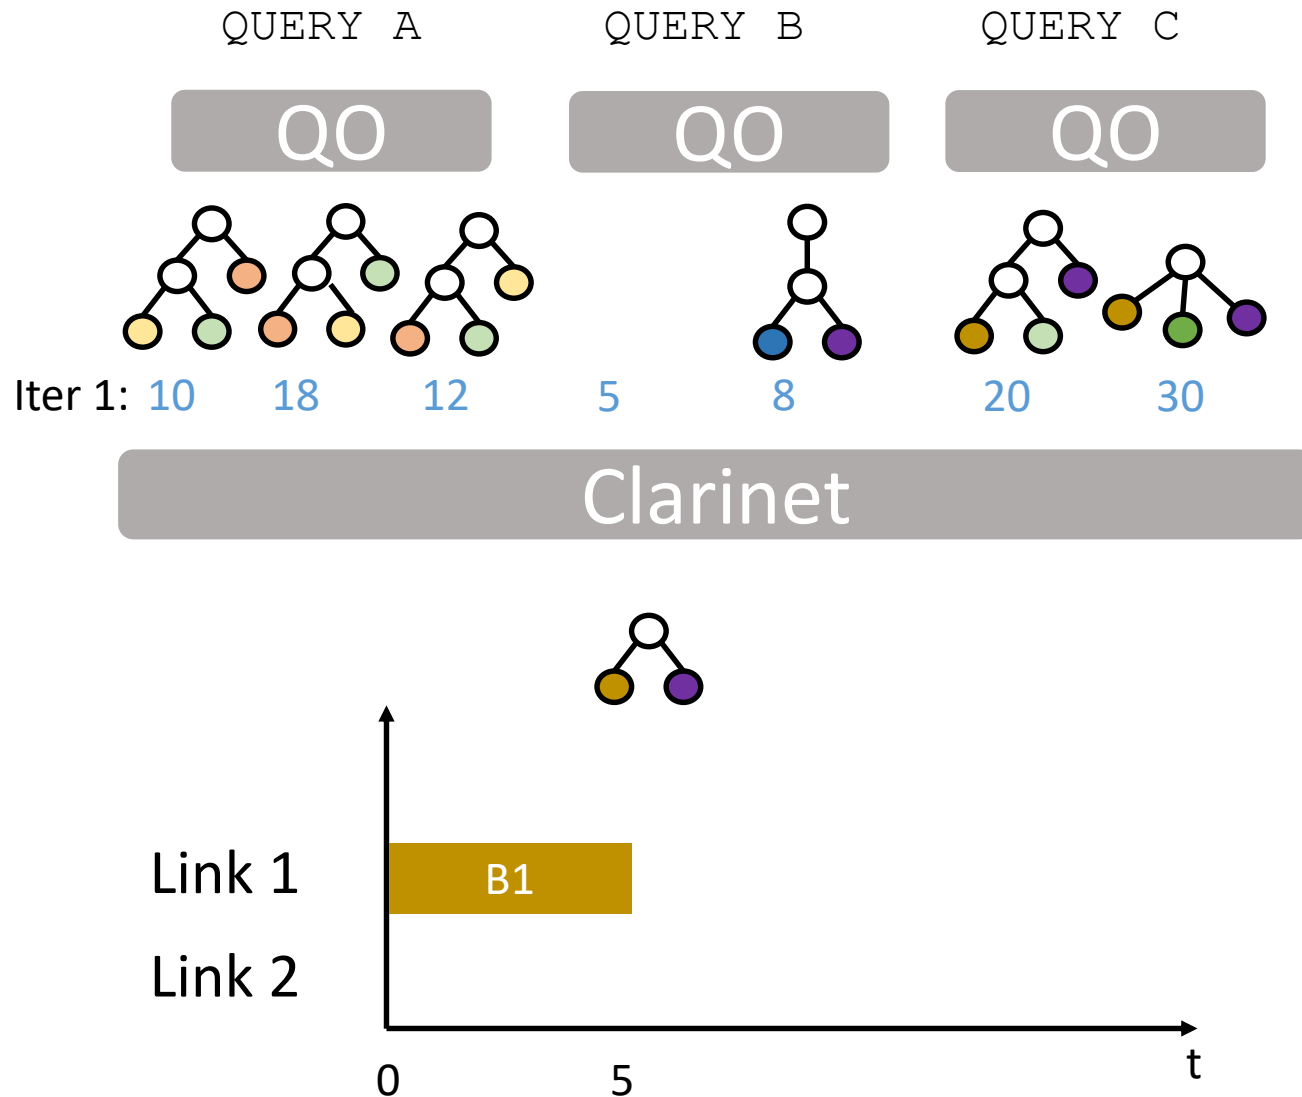
- Best combination → minimize average completion
 - Computationally intractable
- Iterative Shortest Job First (SJF) scheduling heuristic
 1. Pick shortest physical query plan in each iteration

Iterative Shortest Job First




- Best combination → minimize average completion
 - Computationally intractable
- Iterative Shortest Job First (SJF) scheduling heuristic
 1. Pick shortest physical query plan in each iteration

Iterative Shortest Job First



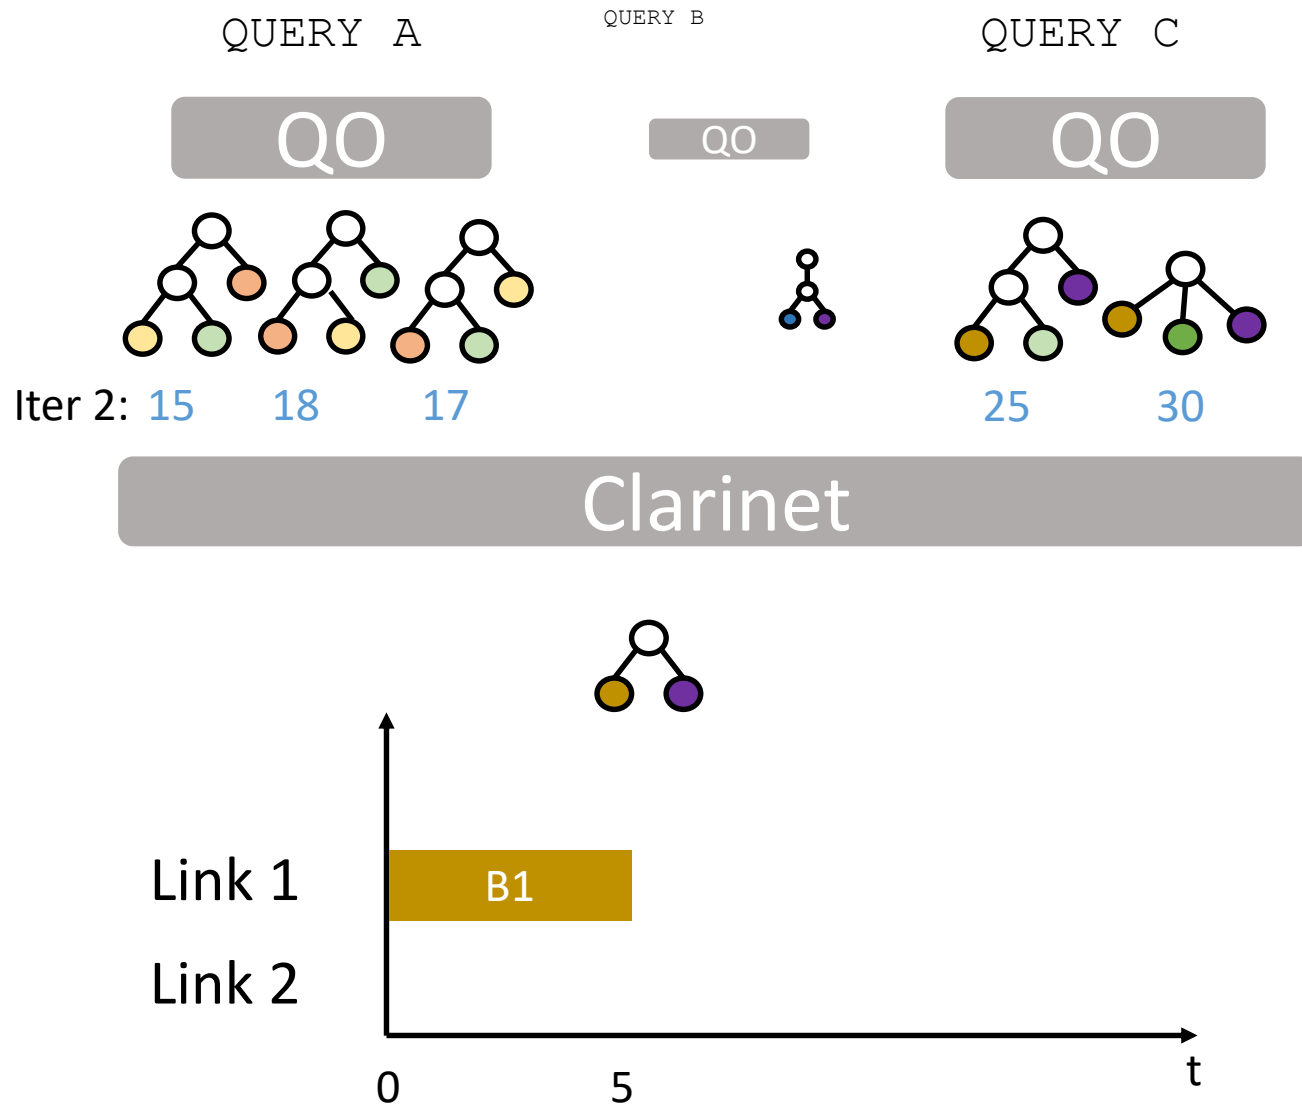
- Best combination → minimize average completion
 - Computationally intractable
- Iterative Shortest Job First (SJF) scheduling heuristic
 1. Pick shortest physical query plan in each iteration
- Reserve bandwidth to guarantee completion time

Iterative Shortest Job First




- Best combination → minimize average completion
 - Computationally intractable
- Iterative Shortest Job First (SJF) scheduling heuristic
 1. Pick shortest physical query plan in each iteration
- Reserve bandwidth to guarantee completion time

Iterative Shortest Job First



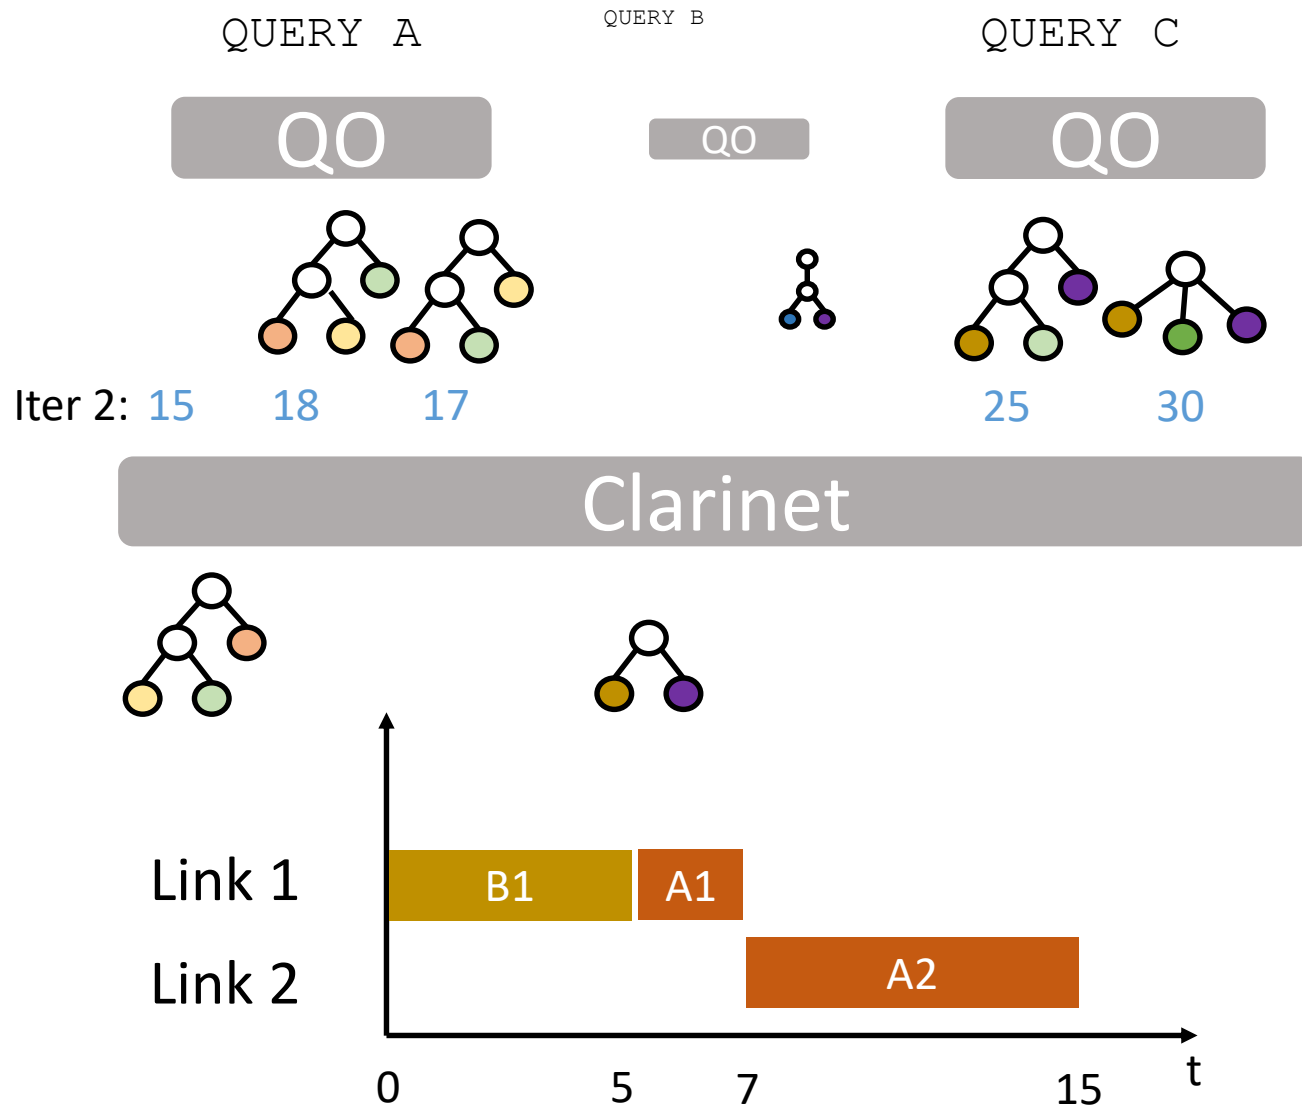
- Best combination → minimize average completion
 - Computationally intractable
- Iterative Shortest Job First (SJF) scheduling heuristic
 1. Pick shortest physical query plan in each iteration
- Reserve bandwidth to guarantee completion time

Iterative Shortest Job First



- Best combination → minimize average completion
 - Computationally intractable
- Iterative Shortest Job First (SJF) scheduling heuristic
 1. Pick shortest physical query plan in each iteration
- Reserve bandwidth to guarantee completion time

Iterative Shortest Job First



- Best combination → minimize average completion
 - Computationally intractable
- Iterative Shortest Job First (SJF) scheduling heuristic
 1. Pick shortest physical query plan in each iteration
- Reserve bandwidth to guarantee completion time


Avoid fragmentation and improve completion time

Avoid fragmentation and improve completion time

- SJF & reservation leads to bandwidth fragmentation


Avoid fragmentation and improve completion time

- SJF & reservation leads to bandwidth fragmentation



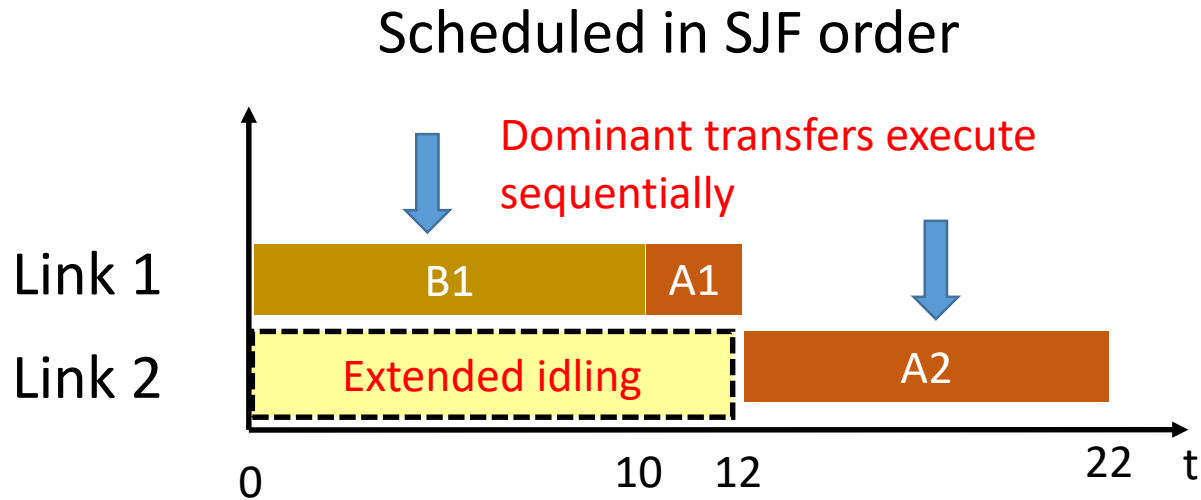
Avoid fragmentation and improve completion time

- SJF & reservation leads to bandwidth fragmentation



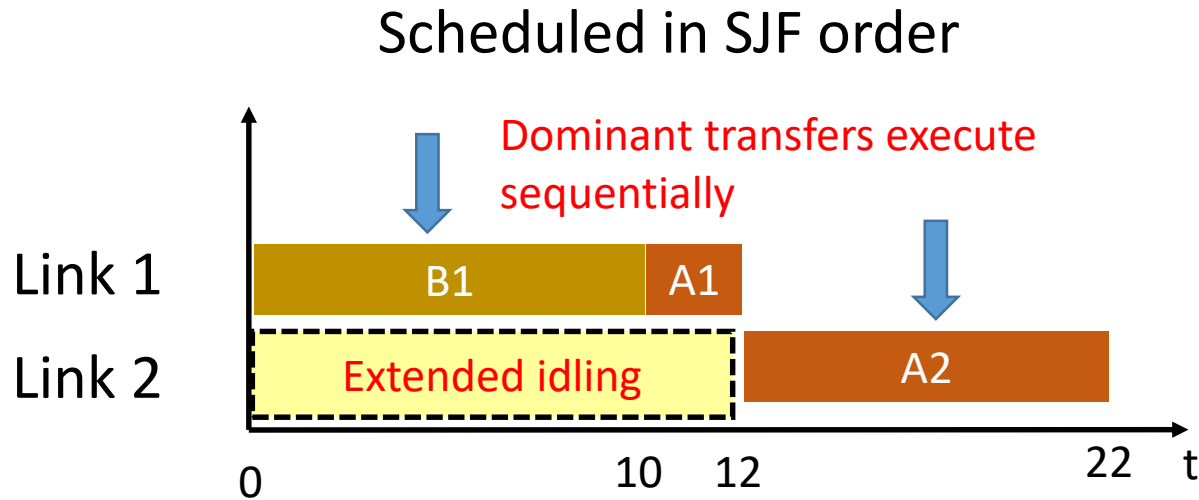
Avoid fragmentation and improve completion time

- SJF & reservation leads to bandwidth fragmentation




Avoid fragmentation and improve completion time

- SJF & reservation leads to bandwidth fragmentation




Alternate schedule with same query plans




Avoid fragmentation and improve completion time

- SJF & reservation leads to bandwidth fragmentation

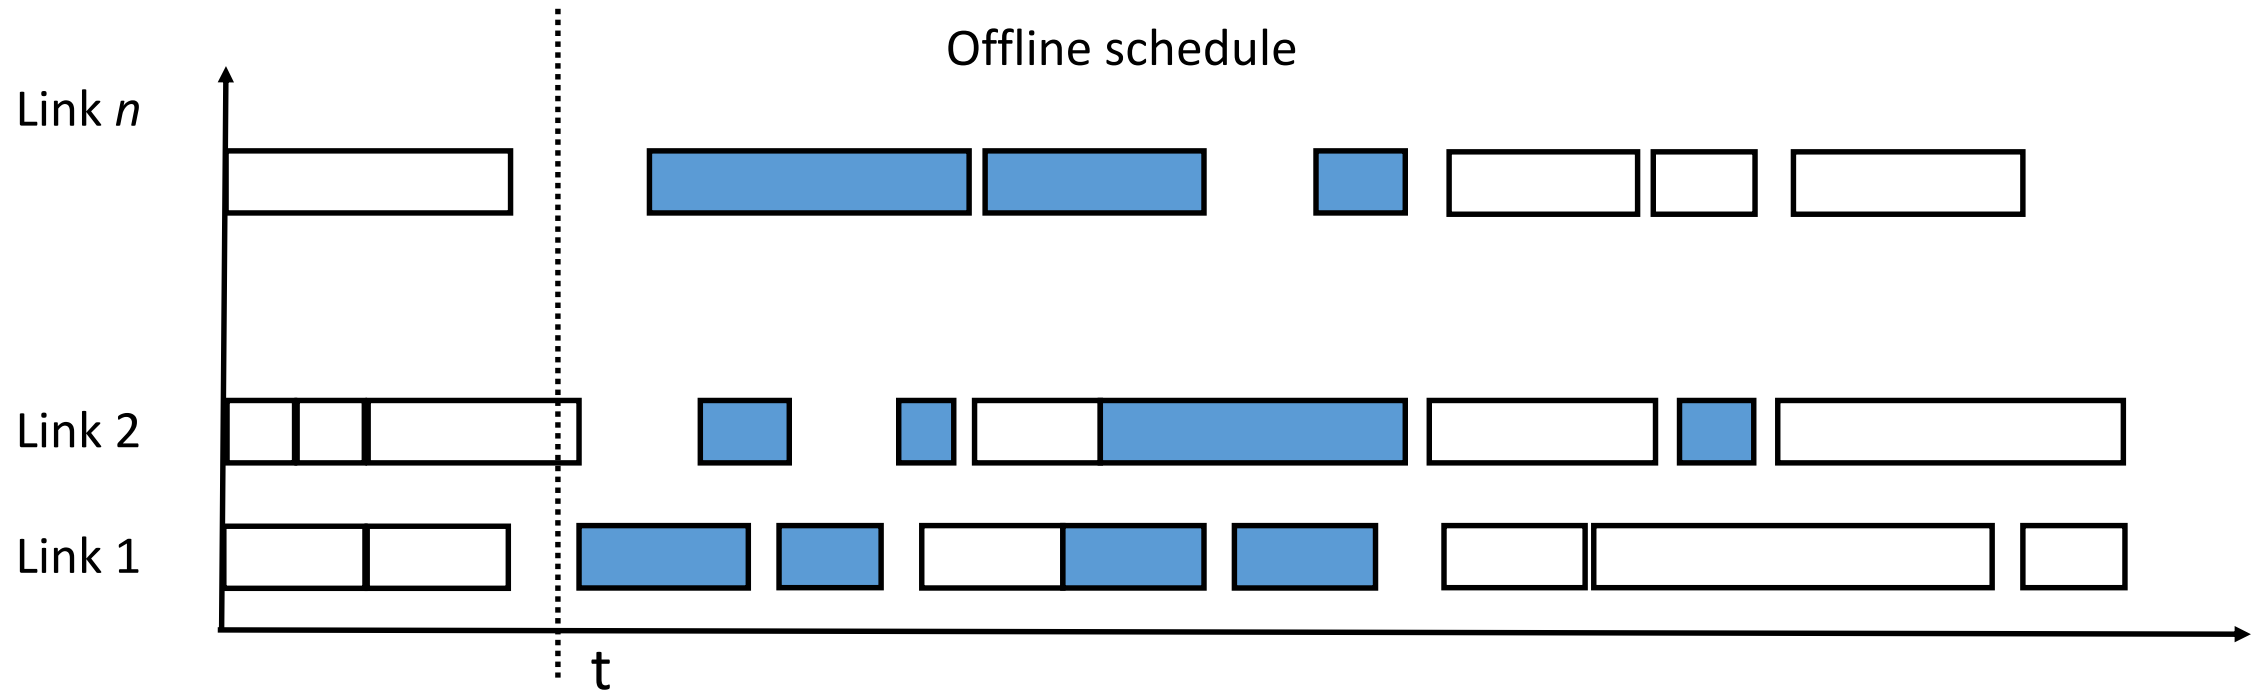


k-Shortest Jobs First Heuristic

Offline schedule




k-Shortest Jobs First Heuristic



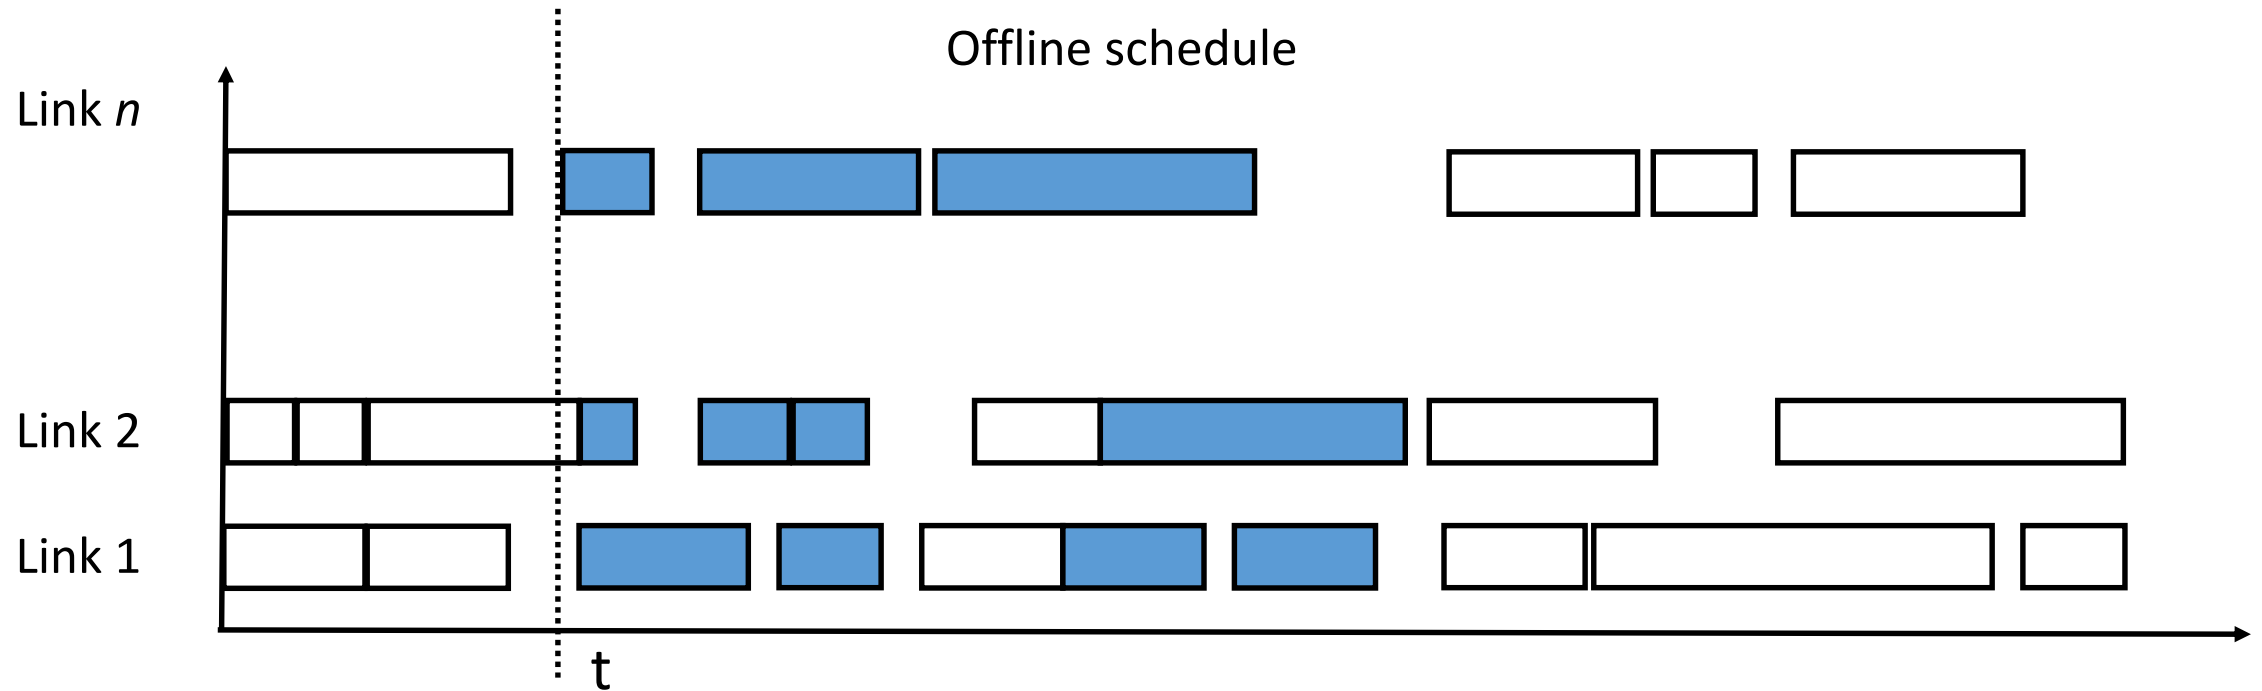
- Identify transfers of k-shortest yet incomplete jobs

k-Shortest Jobs First Heuristic




- Identify transfers of k -shortest yet incomplete jobs
- Relax transfer schedule \rightarrow Start as soon as link is free and task is available

k-Shortest Jobs First Heuristic



- Identify transfers of k-shortest yet incomplete jobs
- Relax transfer schedule → Start as soon as link is free and task is available
- Best 'k' ← Prior observations (or) through offline simulations


Clarinet Implementation



Batch of queries

Existing Query Optimizers

Clarinet Implementation




Batch of queries

Existing Query Optimizers

- Modified Hive to generate multiple plans

Clarinet Implementation




Batch of queries

Existing Query Optimizers

- Modified Hive to generate multiple plans
- QOs control set of generated plans
- Existing optimizations are applied
 - Push down Select
 - Partition pruning

Clarinet Implementation



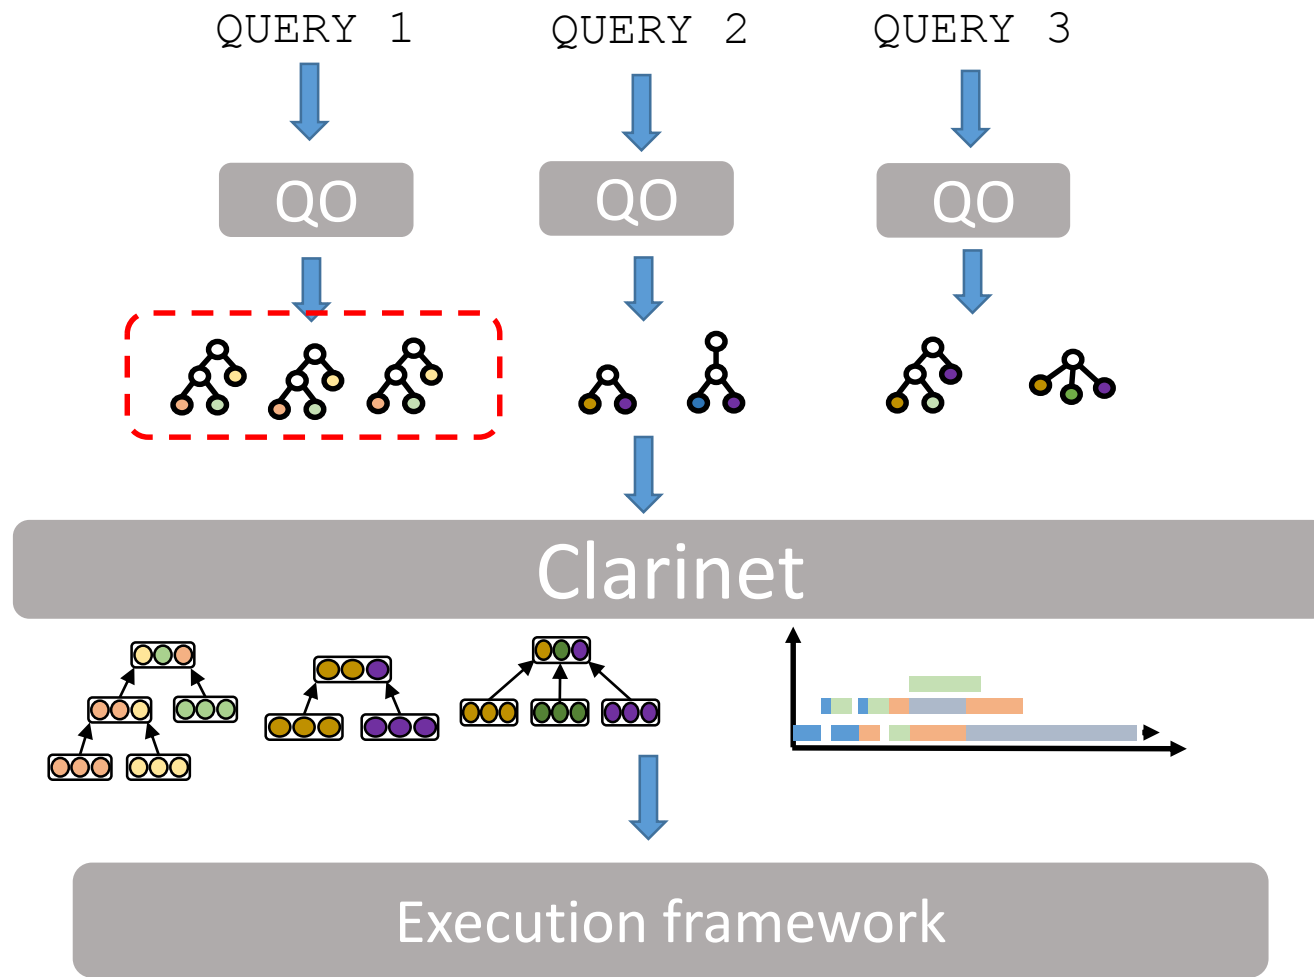
Batch of queries

Existing Query Optimizers

- Modified Hive to generate multiple plans
- QOs control set of generated plans
- Existing optimizations are applied
 - Push down Select
 - Partition pruning

Enforces Clarinet's schedule

Clarinet Implementation



Batch of queries


Existing Query Optimizers

- Modified Hive to generate multiple plans
- QOs control set of generated plans
- Existing optimizations are applied
 - Push down Select
 - Partition pruning

Enforces Clarinet's schedule

- Modified Tez's DAGScheduler

Clarinet Implementation



Batch of queries

Online query arrivals

Existing Query Optimizers

- Modified Hive to generate multiple plans
- QOs control set of generated plans
- Existing optimizations are applied
 - Push down Select
 - Partition pruning

Enforces Clarinet's schedule

- Modified Tez's DAGScheduler
- Fairness guarantees

Evaluation

Compare Clarinet with following GDA approaches:

Evaluation

Compare Clarinet with following GDA approaches:

1. Hive
2. Hive + Iridium
3. Hive + Reducers in single DC

Evaluation

Compare Clarinet with following GDA approaches:

1. Hive : WAN agnostic task placement + scheduling
2. Hive + Iridium
3. Hive + Reducers in single DC

Evaluation

Compare Clarinet with following GDA approaches:

1. Hive : WAN agnostic task placement + scheduling
2. Hive + Iridium : WAN aware task placement across DCs
3. Hive + Reducers in single DC

Evaluation

Compare Clarinet with following GDA approaches:

1. Hive : WAN agnostic task placement + scheduling
2. Hive + Iridium : WAN aware task placement across DCs
3. Hive + Reducers in single DC : Distributed filtering + central aggregation

Evaluation

Compare Clarinet with following GDA approaches:

1. Hive : WAN agnostic task placement + scheduling
 2. Hive + Iridium : WAN aware task placement across DCs
 3. Hive + Reducers in single DC : Distributed filtering + central aggregation
- Geo-Distributed Analytics stack across 10 EC2 regions

Evaluation

Compare Clarinet with following GDA approaches:

1. Hive : WAN agnostic task placement + scheduling
 2. Hive + Iridium : WAN aware task placement across DCs
 3. Hive + Reducers in single DC : Distributed filtering + central aggregation
- Geo-Distributed Analytics stack across 10 EC2 regions
 - Workload:
 - 30 batches of 12 randomly chosen TPC-DS queries

Evaluation: Reduction in average completion time

GDA Approach Vs. Hive	Average Gains
Clarinet	2.7x
Hive + Iridium	1.5x


Evaluation: Reduction in average completion time

GDA Approach Vs. Hive	Average Gains
Clarinet	2.7x
Hive + Iridium	1.5x

Clarinet chooses a different plan
for 75% of queries

Evaluation: Reduction in average completion time

GDA Approach Vs. Hive	Average Gains
Clarinet	2.7x
Hive + Iridium	1.5x




Data from a single batch 12 queries

Clarinet chooses a different plan
for 75% of queries

Evaluation: Reduction in average completion time

GDA Approach Vs. Hive	Average Gains
Clarinet	2.7x
Hive + Iridium	1.5x




Data from a single batch 12 queries

Clarinet chooses a different plan
for 75% of queries

Evaluation: Reduction in average completion time

GDA Approach Vs. Hive	Average Gains
Clarinet	2.7x
Hive + Iridium	1.5x
Hive + Reducers in single DC	0.6x



Data from a single batch 12 queries

Clarinet chooses a different plan
for 75% of queries

Evaluation: Optimization overhead

Evaluation: Optimization overhead

1. Generate multiple query plans
2. Iterative multi-query plan selection

Evaluation: Optimization overhead

1. Generate multiple query plans
 - Up to 64 plans in less than 5 s
2. Iterative multi-query plan selection

Evaluation: Optimization overhead

1. Generate multiple query plans
 - Up to 64 plans in less than 5 s
2. Iterative multi-query plan selection
 - Max. 15 s for batches with 12 queries


Evaluation: Optimization overhead

1. Generate multiple query plans
 - Up to 64 plans in less than 5 s
2. Iterative multi-query plan selection
 - Max. 15 s for batches with 12 queries

Insignificant w.r.t. query running times (order of 10's of minutes)


Summary

- WAN-awareness in QO + cross-layer optimization



Summary

- WAN-awareness in QO + cross-layer optimization
- Presented a **scalable** way to implement **multi-query optimization** with minimal overhead



Summary

- WAN-awareness in QO + cross-layer optimization
- Presented a **scalable** way to implement **multi-query optimization** with minimal overhead

